

BDA HM3

Piero Birello

January 2025

Code: BernMetropolis.R

```
1 graphics.off()
2 rm(list=ls(all=TRUE))
3 fileNameRoot="Figures/BernMetrop" # for output filenames
4 source("DBDA2E-utilities.R")
5
6 # Specify the data, to be used in the likelihood function.
7 myData = c(0,1,1)
8 #myData = c(rep(0,6),rep(1,14))
9
10 # Define the Bernoulli likelihood function, p(D|theta).
11 # The argument theta could be a vector, not just a scalar.
12 likelihood = function( theta , data ) {
13   z = sum( data )
14   N = length( data )
15   pDataGivenTheta = theta^z * (1-theta)^(N-z)
16   # The theta values passed into this function are generated at random,
17   # and therefore might be inadvertently greater than 1 or less than 0.
18   # The likelihood for theta > 1 or for theta < 0 is zero:
19   pDataGivenTheta[ theta > 1 | theta < 0 ] = 0
20   return( pDataGivenTheta )
21 }
22
23 # Define the prior density function.
24 prior = function( theta ) {
25   #pTheta = dbeta( theta , 1 , 1 )
26   pTheta = (cos(4*pi*theta)+1) ** 2/1.5
27   # The theta values passed into this function are generated at random,
28   # and therefore might be inadvertently greater than 1 or less than 0.
29   # The prior for theta > 1 or for theta < 0 is zero:
30   pTheta[ theta > 1 | theta < 0 ] = 0
31   return( pTheta )
32 }
33
34 # Define the relative probability of the target distribution,
35 # as a function of vector theta. For our application, this
36 # target distribution is the unnormalized posterior distribution.
37 targetRelProb = function( theta , data ) {
38   targetRelProb = likelihood( theta , data ) * prior( theta )
39   return( targetRelProb )
40 }
41
42 # Specify the length of the trajectory, i.e., the number of jumps to try:
43 trajLength = 50000 # arbitrary large number
44 # Initialize the vector that will store the results:
45 trajectory = rep( 0 , trajLength )
```

```

46 # Specify where to start the trajectory:
47 trajectory[1] = 0.99 #0.01 # arbitrary value
48 # Specify the burn-in period:
49 burnIn = ceiling( 0.0 * trajLength ) # arbitrary number, less than trajLength
50 # Initialize accepted, rejected counters, just to monitor performance:
51 nAccepted = 0
52 nRejected = 0
53
54 # Now generate the random walk. The 't' index is time or trial in the walk.
55 # Specify seed to reproduce same random walk:
56 set.seed(47406)
57 # Specify standard deviation of proposal distribution:
58 proposalSD = c(0.02,0.2,2.0)[1]
59 for ( t in 1:(trajLength-1) ) {
60     currentPosition = trajectory[t]
61     # Use the proposal distribution to generate a proposed jump.
62     proposedJump = rnorm( 1 , mean=0 , sd=proposalSD )
63     # Compute the probability of accepting the proposed jump.
64     probaAccept = min( 1,
65         targetRelProb( currentPosition + proposedJump , myData )
66         / targetRelProb( currentPosition , myData ) )
67     # Generate a random uniform value from the interval [0,1] to
68     # decide whether or not to accept the proposed jump.
69     if ( runif(1) < probaAccept ) {
70         # accept the proposed jump
71         trajectory[ t+1 ] = currentPosition + proposedJump
72         # increment the accepted counter, just to monitor performance
73         if ( t > burnIn ) { nAccepted = nAccepted + 1 }
74     } else {
75         # reject the proposed jump, stay at current position
76         trajectory[ t+1 ] = currentPosition
77         # increment the rejected counter, just to monitor performance
78         if ( t > burnIn ) { nRejected = nRejected + 1 }
79     }
80 }
81
82 # Extract the post-burnIn portion of the trajectory.
83 acceptedTraj = trajectory[ (burnIn+1) : length(trajectory) ]
84
85 # End of Metropolis algorithm.
86
87 #-----
88 # Display the chain.
89
90 openGraph(width=4,height=8)
91 layout( matrix(1:3,nrow=3) )
92 par(mar=c(3,4,2,1),mgp=c(2,0.7,0))
93
94 # Posterior histogram:
95 paramInfo = plotPost( acceptedTraj , xlim=c(0,1) , xlab=bquote(theta) ,
96     cex.main=2.0 ,
97     main=bquote( list( "Prpsl.SD" == .(proposalSD) ,
98         "Eff.Sz." == .(round(effectiveSize(acceptedTraj),1)) ) ) )
99
100 # Trajectory, a.k.a. trace plot, end of chain:
101 idxToPlot = (trajLength-100):trajLength
102 plot( trajectory[idxToPlot] , idxToPlot , main="End of Chain" ,
103     xlab=bquote(theta) , xlim=c(0,1) , ylab="Step in Chain" ,
104     type="o" , pch=20 , col="skyblue" , cex.lab=1.5 )
105 # Display proposal SD and acceptance ratio in the plot.
106 text( 0.0 , trajLength , adj=c(0.0,1.1) , cex=1.75 ,
107     labels = bquote( frac(N[acc],N[pro]) ==

```

```

108         .(signif( nAccepted/length(acceptedTraj) , 3 )))
109
110 # Trajectory, a.k.a. trace plot, beginning of chain:
111 idxToPlot = 1:100
112 plot( trajectory[idxToPlot] , idxToPlot , main="Beginning of Chain" ,
113       xlab=bquote(theta) , xlim=c(0,1) , ylab="Step in Chain" ,
114       type="o" , pch=20 , col="skyblue" , cex.lab=1.5 )
115 # Indicate burn in limit (might not be visible if not in range):
116 if ( burnIn > 0 ) {
117     abline(h=burnIn,lty="dotted")
118     text( 0.5 , burnIn+1 , "Burn In" , adj=c(0.5,1.1) )
119 }
120
121 saveGraph( file=paste0( fileNameRoot ,
122                        "SD" , proposalSD ,
123                        "Init" , trajectory[1] , "Posterior" ) , type="eps" )
124
125 # Open a plot, specifying height and width
126 openGraph(height=7,width=3.5)
127 # Create 2-rows layout for two plots
128 layout(matrix(1:2,nrow=2))
129 # Plot autocorrelation function of the accepted traj
130 acf( acceptedTraj , lag.max=30 , col="skyblue" , lwd=3 )
131 # Get accepted traj length
132 Len = length( acceptedTraj )
133 # Define lag
134 Lag = 10
135 # Define head of the traj, from 1 to Len-Lag
136 trajHead = acceptedTraj[ 1
137                        : (Len-Lag) ]
138 # Define tail of the traj, from 1+Lag to Len
139 trajTail = acceptedTraj[ (1+Lag) : Len
140                       ]
141 # Scatterplot of the lagged trajectories
142 plot( trajHead , trajTail , pch="." , col="skyblue" ,
143       main=bquote( list( "Prpsl.SD" == .(proposalSD) ,
144                         lag == .(Lag) ,
145                         cor == .(round(cor(trajHead,trajTail),3)))) )
146
147 saveGraph( file=paste0( fileNameRoot ,
148                        "SD" , proposalSD ,
149                        "Init" , trajectory[1] , "ACF" , "Posterior" ) , type="eps" )
150
151 #-----

```

7.1

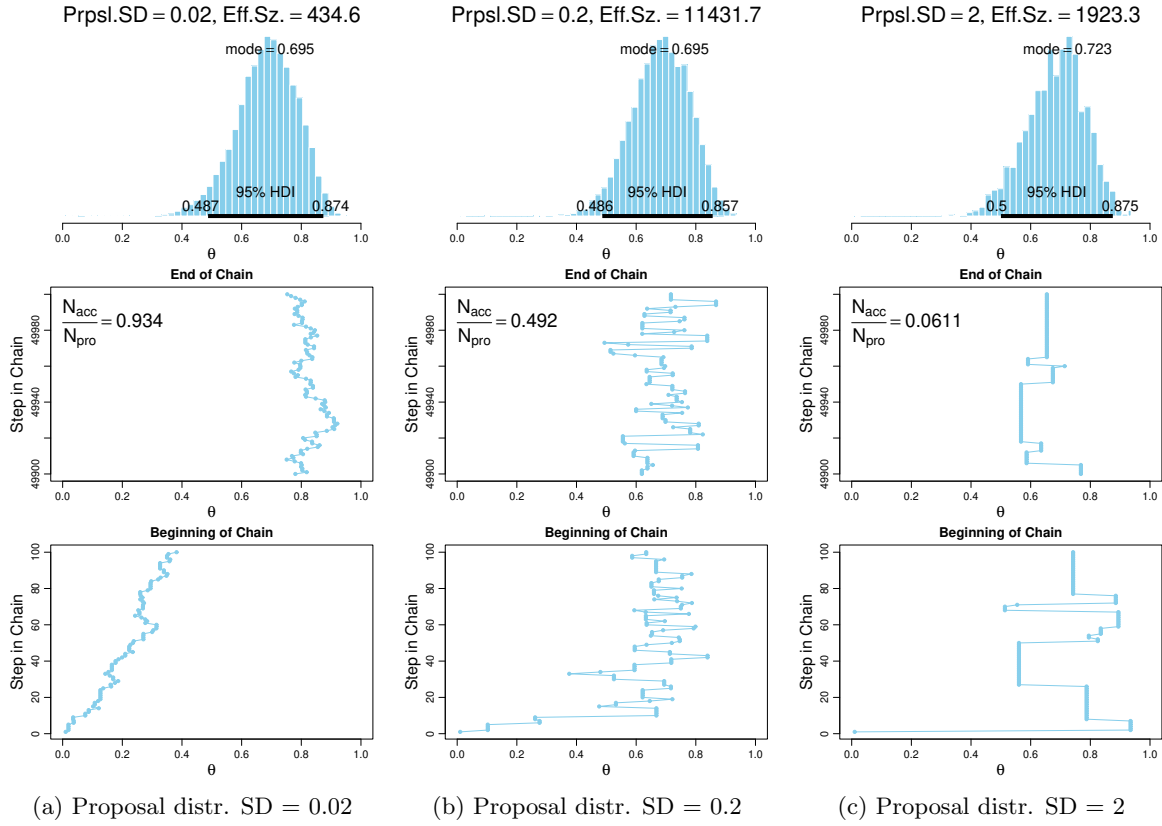


Figure 1: Repetition of the experiment of Figure 7.4 of the book, with a different random seed. The same considerations hold, i.e., the most efficient sampling is the one with proposal SD = 0.2. When SD = 0.02, steps in the chain are too small and almost always accepted, hence the algorithm will take long time to sample the target distribution. As a result, the ESS is low. The opposite holds for SD = 2, where steps in the chain are too large, so that the acceptance ratio is very low. Again, this leads to low ESS.

7.2

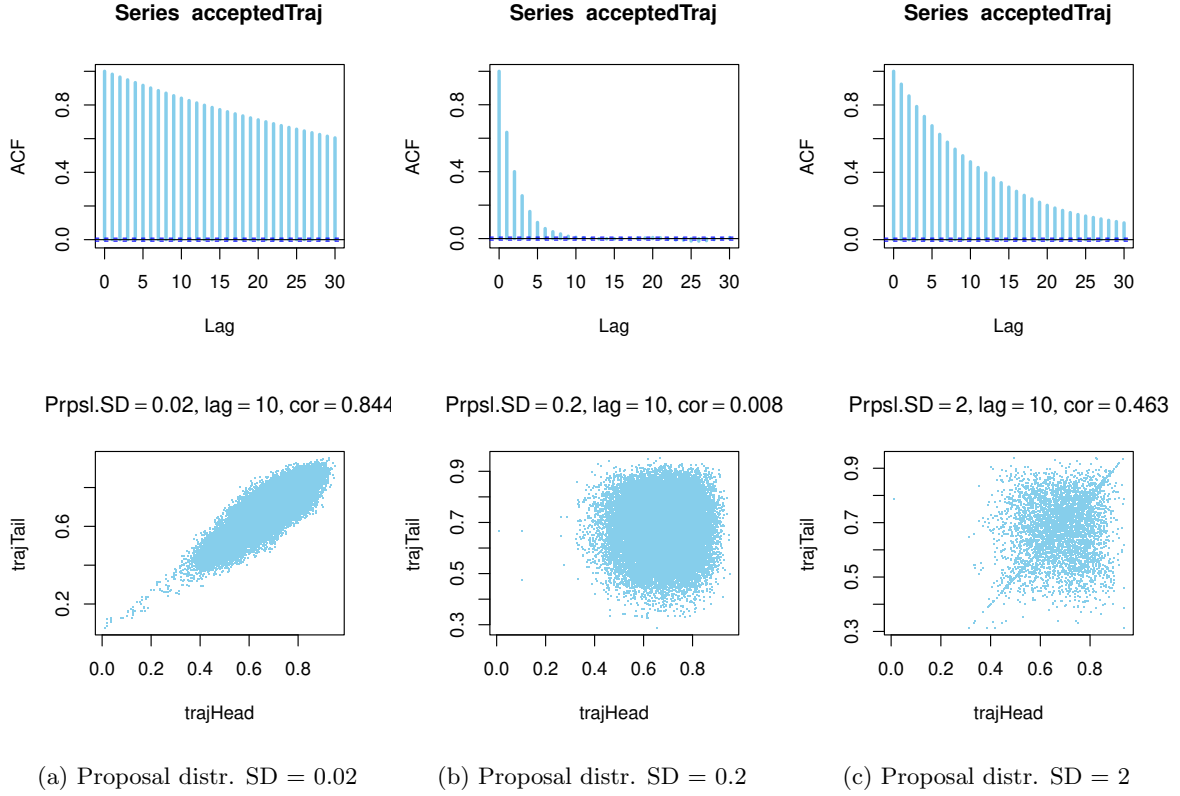


Figure 2: Here we show the autocorrelation function (ACF) for the trajectories described in 7.1. In the top row, ACF is plotted as a function of the lag time. In the bottom row, we show two copies of the trajectories with a lag of 10 steps, and compute the correlation for this single value. Plots in the top and bottom row are consistent, i.e., ACF values match. The trajectory resulting from $\text{SD} = 0.2$ is the one showing the fastest decaying ACF. Note that when $\text{SD} = 2$, we observe a dense line of points on the diagonal of the scatter plot in the bottom row. Actually, being the acceptance ratio very low, the trajectory often gets stuck at some parameter value, even for long periods. When the trajectory is stuck for a number of time steps greater than the lag time (here 10), then the lagged trajectory superimposes with the non lagged one.

7.3

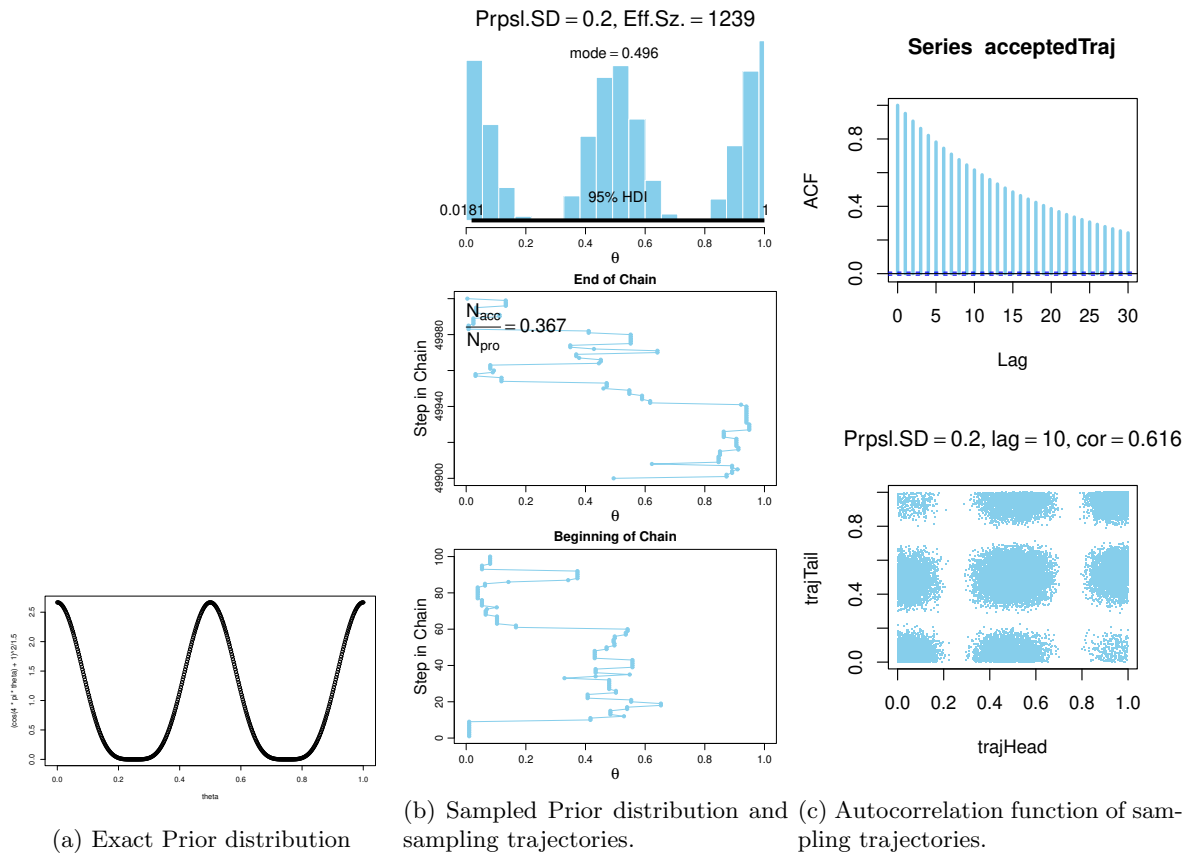
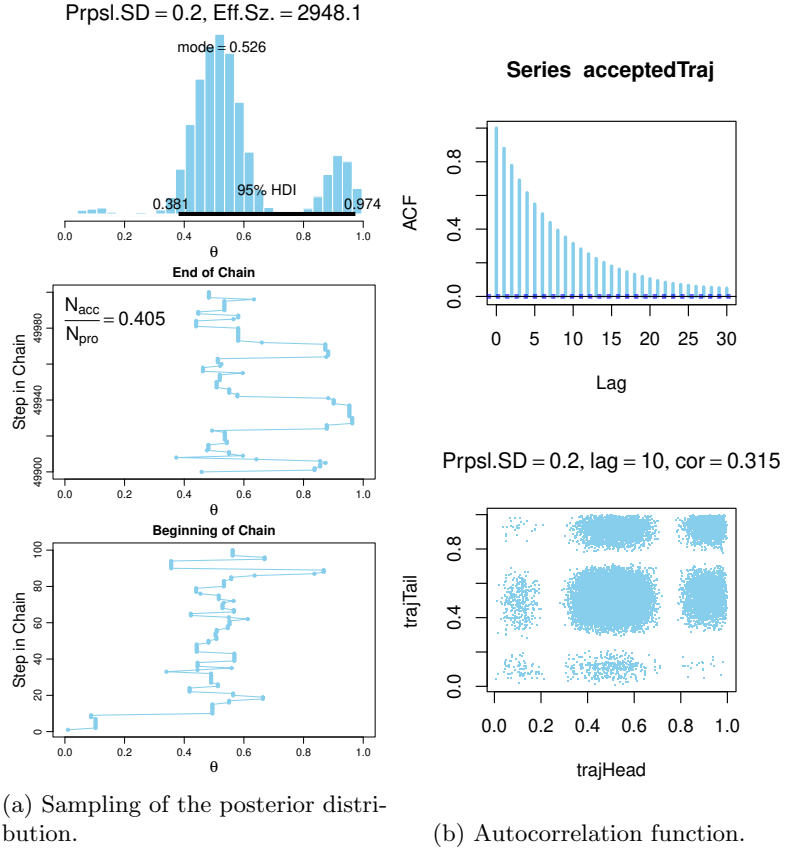


Figure 3: Exact and sampled prior. The sampling is good enough to reproduce the correct shape of the distribution, even though a larger sampling size would smoothen our estimate.



(a) Sampling of the posterior distribution.

(b) Autocorrelation function.

Figure 4: Here we evaluate the posterior distribution given outcomes 0,1,1. Our prior distribution is a three-modal distribution with peaks located at 0, 0.5 and 1 respectively. Our data suggests that, among the three corresponding 'hills', the one on the left is the least likely and the one in the middle is the most likely.

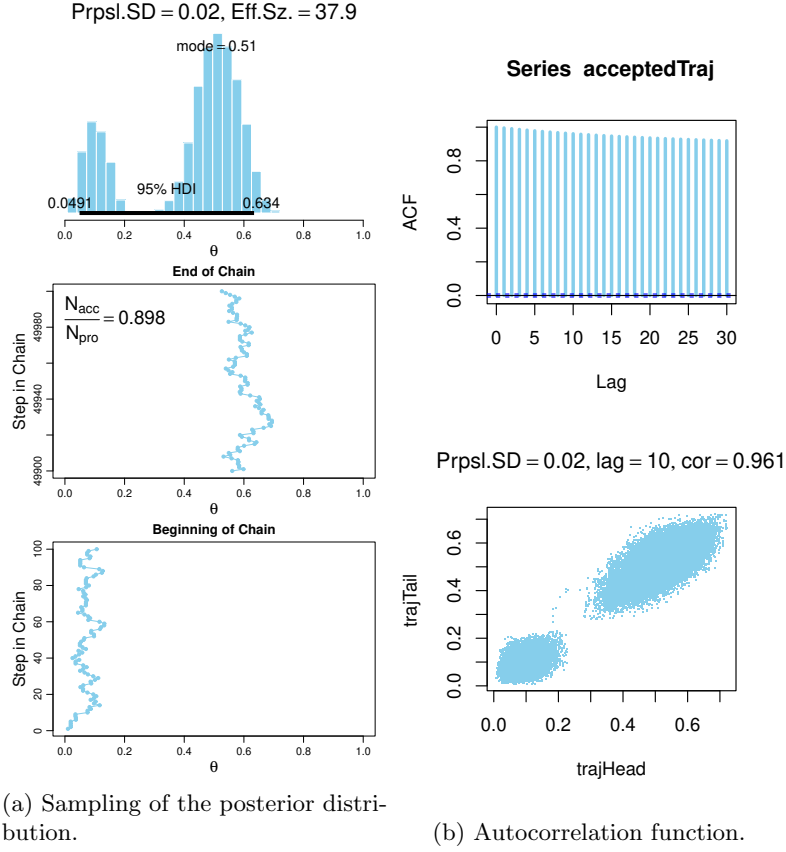
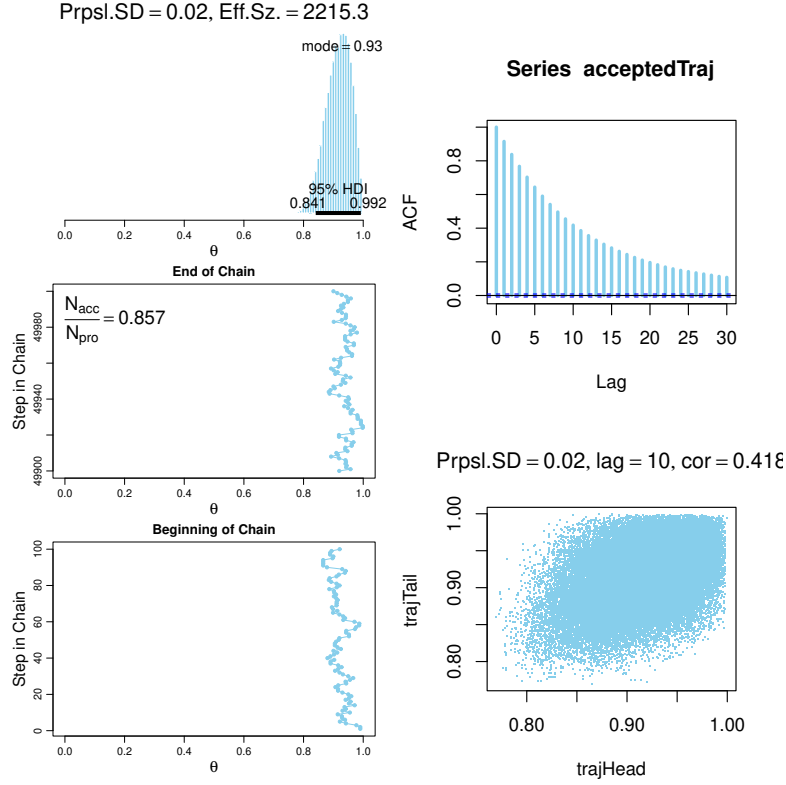


Figure 5: Here we replicate the experiment of the previous figure, but the chosen proposal SD is too small. We are not correctly sampling the posterior distribution, and we get a wrong estimate of it. There are many ways to realize something is going wrong, including a very small ESS and very slowly decaying ACF. Also, if you look at the next figure...



(a) Sampling of the posterior distribution.

(b) Autocorrelation function.

Figure 6: ...changing the initial condition we get a completely different estimate! Also notice that in this case the random walk remains stuck in the right peak, and is so unaware of the unobserved part of the distribution that ACF behaves well, and the effective size is consequently large. Trying different initial conditions turns out to be crucial.